



How many ReLU random features before an RF denoiser may beat the linear one?


MNIST ($d = 784$): Wiener and empirical-prior Bayes losses



CIFAR-10 ($d = 3072$): Wiener and empirical-prior Bayes losses



MNIST: widths k at which the gain is guaranteed small



CIFAR-10: widths k at which the gain is guaranteed small

