



Held out: pool/fit = images [0:10000], test = [10000:10500], k and τ picked on the last 100 images. Networks: hollow = $\sigma < 0.001$ (below training range). train = eval: pool images [0:500], pool contains each.



- Linear (Wiener) — 10k pool cov (rank 9999)
- Linear (Wiener) — larger full-rank cov
- Softmax, global — held out
- Softmax, local — Lukoianov (Wiener mask)
- Lukoianov — mask from larger covariance
- Softmax, local — Kamb & Ganguli LS
- U-net, all 15.8k, long (test seen)
- EDM U-net (official, all 15.8k; test seen)
- Softmax, global — train = eval

Lukoianov at $\sigma \leq 0.001 \rightarrow 0$ is an 8-bit artefact: as $\sigma \rightarrow 0$ its Wiener mask shrinks to the diagonal, a per-coordinate nearest value among 10k training pixels, which snaps to the exact 1/255 level once $\sigma \ll 1/255$. val-selected, by σ ascending: LS k = 1, 1, 1, 1, 1, 1, 1, 1, 3, 7, 7, 11, 11, 15, 23, 39, 63, 63, 39; Lukoianov τ = 0.3, 0.5, 0.5, 0.5, 1, 0.3, 0.3, 0.2, 0.2, 0.1, 0.1, 0.1, 0.05, 0.05, 0.05, 0.05, 0.05, 0.1, 0.2