



Held out: pool/fit = images [0:10000], test = [10000:10500], k and τ picked on the last 100 images. Networks: hollow = $\sigma < 0.001$ (below training range). train = eval: pool images [0:500], pool contains each.

MSE vs noise, AFHQ 64×64



Relative to linear (< 1 beats Wiener)



Lukoianov at $\sigma \leq 0.001 \rightarrow 0$ is an 8-bit artifact: as $\sigma \rightarrow 0$ its Wiener mask shrinks to the diagonal, a per-coordinate nearest value among 10k training pixels, which snaps to the exact 1/255 level once $\sigma \ll 1/255$. val-selected, by σ ascending: LS k = 1, 1, 1, 1, 1, 1, 1, 1, 3, 7, 7, 11, 11, 15, 23, 39, 63, 63, 39; Lukoianov τ = 0.3, 0.5, 0.5, 0.5, 1, 0.3, 0.3, 0.2, 0.2, 0.1, 0.1, 0.1, 0.05, 0.05, 0.05, 0.05, 0.05, 0.1, 0.2