



Held out: pool/fit = images [0:10000], test = [60000:60500], k and τ picked on the last 100 images. Networks: hollow = $\sigma < 0.001$ (below training range). train = eval: pool images [0:500], pool contains each.



- Linear (Wiener) — 10k pool cov (rank 9999)
- Linear (Wiener) — larger full-rank cov
- Softmax, global — held out
- Softmax, local — Lukoianov (Wiener mask)
- Softmax, global — train = eval
- Lukoianov — mask from larger covariance
- Softmax, local — Kamb & Ganguli LS
- U-net, 30k [0:30k], held out
- U-net, 10k = the pool, held out
- EDM U-net (official, all 70k; test seen)

Lukoianov at $\sigma \leq 0.001 \rightarrow 0$ is an 8-bit artefact: as $\sigma \rightarrow 0$ its Wiener mask shrinks to the diagonal, a per-coordinate nearest value among 10k training pixels, which snaps to the exact 1/255 level once $\sigma \ll 1/255$. val-selected, by σ ascending: LS k = 1, 1, 1, 1, 1, 1, 1, 1, 3, 7, 9, 11, 15, 19, 23, 39, 63, 63, 63; Lukoianov τ = 0.5, 0.5, 0.5, 0.5, 0.3, 0.3, 0.3, 0.3, 0.2, 0.1, 0.1, 0.1, 0.05, 0.05, 0.05, 0.05, 0.05, 0.02, 0.02