



Held out: pool/fit = images [0:10000], test = [60000:60500], k and τ picked on the last 100 images. Networks: hollow = $\sigma < 0.001$ (below training range). train = eval: pool images [0:500], pool contains each.

MSE vs noise, FFHQ 64×64



Relative to linear (< 1 beats Wiener)



Lukoianov at $\sigma \leq 0.001 \rightarrow 0$ is an 8-bit artifact: as $\sigma \rightarrow 0$ its Wiener mask shrinks to the diagonal, a per-coordinate nearest value among 10k training pixels, which snaps to the exact 1/255 level once $\sigma \ll 1/255$.
val-selected, by σ ascending: LS k = 1, 1, 1, 1, 1, 1, 1, 1, 3, 7, 9, 11, 15, 19, 23, 39, 63, 63, 63; Lukoianov τ = 0.5, 0.5, 0.5, 0.5, 0.3, 0.3, 0.3, 0.3, 0.2, 0.1, 0.1, 0.1, 0.05, 0.05, 0.05, 0.05, 0.05, 0.02, 0.02